

Training code

STEP1: WE ARE IMPLEMENTING YOLO IN KERAS TO TRY TO UNDERSTAND WHAT WE ARE DOING

Because our business case is quite unique, we might not need all the layers that the original yolo model offers

In [3]:

```
#Let's first load all the libraries we need

from keras.models import Sequential, Model
from keras.layers import Reshape, Activation, Conv2D, Input, MaxPooling2D,
BatchNormalization, Flatten, Dense, Lambda, Dropout
from keras.layers.advanced_activations import LeakyReLU
from keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard
from keras.optimizers import SGD, Adam, RMSprop, Adamax
from keras.layers.merge import concatenate
import matplotlib.pyplot as plt
import keras.backend as K
import tensorflow as tf
#import imgaug as ia
from tqdm import tqdm
#from imgaug import augmenters as iaa
import numpy as np
import pandas as pd
import pickle
import os, cv2
#from preprocessing import parse_annotation, BatchGenerator

/anaconda/envs/py35/lib/python3.5/site-packages/h5py/__init__.py:36: Future
Warning: Conversion of the second argument of issubdtype from `float` to `n
p.floating` is deprecated. In future, it will be treated as `np.float64 ==
np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

In [4]:

```
#custom to us are the labels and the image size
LABELS = ['melanoma', 'notmelanoma']

#I originally tried 200 * 200 and gave me an error - this is because
the input layer is a 32 neuron so we need multiples of 32
#I'm going to use 32*7 = 224

IMAGE_H, IMAGE_W = 192, 192
#Grids are used when you are facing problems with more than one object to
detect and the fact they allow 2 (in the original)
#overlapping bounding boxes. In our case, we have only 1 very well defined
object to detect so we don't need more than 1 grid

GRID_H, GRID_W = 7, 7

#Let's leave the rest as is

BOX = 5
CLASS = len(LABELS)
CLASS_WEIGHTS = np.ones(CLASS, dtype='float32')
```

```

OBJ_THRESHOLD      = 0.3#0.5
NMS_THRESHOLD      = 0.3#0.45
ANCHORS            = [0.57273, 0.677385, 1.87446, 2.06253, 3.33843, 5.47434,
7.88282, 3.52778, 9.77052, 9.16828]

NO_OBJECT_SCALE    = 1.0
OBJECT_SCALE       = 5.0
COORD_SCALE        = 1.0
CLASS_SCALE        = 1.0

BATCH_SIZE         = 16
WARM_UP_BATCHES    = 0
TRUE_BOX_BUFFER    = 50

```

STEP 1.1: LET'S BUILD THE NETWORK

In [6]:

```

# the function to implement the orgnization layer (thanks to
github.com/allanzelener/YAD2K)
def space_to_depth_x2(x):
    return tf.space_to_depth(x, block_size=2)

```

In [8]:

```

input_image = Input(shape=(IMAGE_H, IMAGE_W, 3))
#true_boxes = Input(shape=(1, 1, 1, TRUE_BOX_BUFFER , 4))

#NOTE ON THE SINTAX = This isn't using Sequential(), it's building a
pipeline of layers applied to
#the input_image. So what this is doing x = fn...(f1(x)). Nested functions.
This is very useful when you need to do skip
#connections or you need to do something a bit more complex to the output,
it isn't a clear sequence

# Layer 1
x = Conv2D(32, (3,3), strides=(1,1), padding='same', name='conv_1',
use_bias=False)(input_image)
x = BatchNormalization(name='norm_1')(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Layer 2
x = Conv2D(64, (3,3), strides=(1,1), padding='same', name='conv_2',
use_bias=False)(x)
x = BatchNormalization(name='norm_2')(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Layer 3
x = Conv2D(128, (3,3), strides=(1,1), padding='same', name='conv_3',
use_bias=False)(x)
x = BatchNormalization(name='norm_3')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 4
x = Conv2D(64, (1,1), strides=(1,1), padding='same', name='conv_4',
use_bias=False)(x)
x = BatchNormalization(name='norm_4')(x)
x = LeakyReLU(alpha=0.1)(x)

```

```

# Layer 5
x = Conv2D(128, (3,3), strides=(1,1), padding='same', name='conv_5',
use_bias=False)(x)
x = BatchNormalization(name='norm_5')(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Layer 6
x = Conv2D(256, (3,3), strides=(1,1), padding='same', name='conv_6',
use_bias=False)(x)
x = BatchNormalization(name='norm_6')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 7
x = Conv2D(128, (1,1), strides=(1,1), padding='same', name='conv_7',
use_bias=False)(x)
x = BatchNormalization(name='norm_7')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 8
x = Conv2D(256, (3,3), strides=(1,1), padding='same', name='conv_8',
use_bias=False)(x)
x = BatchNormalization(name='norm_8')(x)
x = LeakyReLU(alpha=0.1)(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Layer 9
x = Conv2D(512, (3,3), strides=(1,1), padding='same', name='conv_9',
use_bias=False)(x)
x = BatchNormalization(name='norm_9')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 10
x = Conv2D(256, (1,1), strides=(1,1), padding='same', name='conv_10',
use_bias=False)(x)
x = BatchNormalization(name='norm_10')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 11
x = Conv2D(512, (3,3), strides=(1,1), padding='same', name='conv_11',
use_bias=False)(x)
x = BatchNormalization(name='norm_11')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 12
x = Conv2D(256, (1,1), strides=(1,1), padding='same', name='conv_12',
use_bias=False)(x)
x = BatchNormalization(name='norm_12')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 13
x = Conv2D(512, (3,3), strides=(1,1), padding='same', name='conv_13',
use_bias=False)(x)
x = BatchNormalization(name='norm_13')(x)
x = LeakyReLU(alpha=0.1)(x)

```

```

skip_connection = x

x = MaxPooling2D(pool_size=(2, 2))(x)

# Layer 14
x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_14',
use_bias=False)(x)
x = BatchNormalization(name='norm_14')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 15
x = Conv2D(512, (1,1), strides=(1,1), padding='same', name='conv_15',
use_bias=False)(x)
x = BatchNormalization(name='norm_15')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 16
x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_16',
use_bias=False)(x)
x = BatchNormalization(name='norm_16')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 17
x = Conv2D(512, (1,1), strides=(1,1), padding='same', name='conv_17',
use_bias=False)(x)
x = BatchNormalization(name='norm_17')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 18
x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_18',
use_bias=False)(x)
x = BatchNormalization(name='norm_18')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 19
x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_19',
use_bias=False)(x)
x = BatchNormalization(name='norm_19')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 20
x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_20',
use_bias=False)(x)
x = BatchNormalization(name='norm_20')(x)
x = LeakyReLU(alpha=0.1)(x)

# Layer 21
skip_connection = Conv2D(64, (1,1), strides=(1,1), padding='same',
name='conv_21', use_bias=False)(skip_connection)
skip_connection = BatchNormalization(name='norm_21')(skip_connection)
skip_connection = LeakyReLU(alpha=0.1)(skip_connection)
skip_connection = Lambda(space_to_depth_x2)(skip_connection)

x = concatenate([skip_connection, x])

# Layer 22

```

```

x = Conv2D(1024, (3,3), strides=(1,1), padding='same', name='conv_22',
use_bias=False)(x)
x = BatchNormalization(name='norm_22')(x)
x = LeakyReLU(alpha=0.1)(x)
# As per the comment below, we only need 4 outputs so we are adding a dense
layer as an output with one neuron per dimension
#Important - the output of layer22 is a 6*6 feature map with 1024 feature
maps so 6*6*1024 - This is not really what we need
#Therefore we need to flatten the output so we can apply the final
transforamtion in the dense and get the output we nee
#output.shape(none, 4)
#none = the number of batches if we were splitting our dataset in batches

x = Flatten()(x)
x = Dropout(0.5)(x)
output = Dense(4, activation='linear')(x)

# Layer 23
#What they are doing here is to recreante a dense layer using a convolusion
but they are just fixing the weights for each feature
#map but reading each pixel at a time with those same weights in each
feature map
#What we want is to adapt this layer to the number of outputs we want. In
our case - xmin, ymin, xmax, ymax

#x = Conv2D(BOX * (4 + 1 + CLASS), (1,1), strides=(1,1), padding='same',
name='conv_23')(x)
#output = Reshape((GRID_H, GRID_W, BOX, 4 + 1 + CLASS))(x)

# small hack to allow true_boxes to be registered when Keras build the
model
# for more information: https://github.com/fchollet/keras/issues/2790
#output = Lambda(lambda args: args[0])([output, true_boxes])

model = Model(input_image, output)

*Freeze layers*

```

Freezing layers - this are needed because when we printed the original summary we have 50million+ trainable parameters and even if we have prelearnt weights, it will optimise them again

Total params: 50,695,396 Trainable params: 50,674,724 Non-trainable params: 20,672

In []:

```

#for i in range(1,23):
    #model.layers[i].trainable = False

```

In [9]:

```
model.summary()
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_3 (InputLayer)	(None, 192, 192, 3)	0	

conv_1 (Conv2D)	(None, 192, 192, 32)	864	input_3[0]
norm_1 (BatchNormalization)	(None, 192, 192, 32)	128	conv_1[0][0]
leaky_re_lu_44 (LeakyReLU)	(None, 192, 192, 32)	0	norm_1[0][0]
max_pooling2d_11 (MaxPooling2D)	(None, 96, 96, 32)	0	leaky_re_lu_44[0][0]
conv_2 (Conv2D)	(None, 96, 96, 64)	18432	max_pooling2d_11[0][0]
norm_2 (BatchNormalization)	(None, 96, 96, 64)	256	conv_2[0][0]
leaky_re_lu_45 (LeakyReLU)	(None, 96, 96, 64)	0	norm_2[0][0]
max_pooling2d_12 (MaxPooling2D)	(None, 48, 48, 64)	0	leaky_re_lu_45[0][0]
conv_3 (Conv2D)	(None, 48, 48, 128)	73728	max_pooling2d_12[0][0]
norm_3 (BatchNormalization)	(None, 48, 48, 128)	512	conv_3[0][0]
leaky_re_lu_46 (LeakyReLU)	(None, 48, 48, 128)	0	norm_3[0][0]
conv_4 (Conv2D)	(None, 48, 48, 64)	8192	leaky_re_lu_46[0][0]
norm_4 (BatchNormalization)	(None, 48, 48, 64)	256	conv_4[0][0]
leaky_re_lu_47 (LeakyReLU)	(None, 48, 48, 64)	0	norm_4[0][0]

conv_5 (Conv2D) u_47[0][0]	(None, 48, 48, 128)	73728	leaky_re_l
norm_5 (BatchNormalization) 0]	(None, 48, 48, 128)	512	conv_5[0][
leaky_re_lu_48 (LeakyReLU) 0]	(None, 48, 48, 128)	0	norm_5[0][
max_pooling2d_13 (MaxPooling2D) u_48[0][0]	(None, 24, 24, 128)	0	leaky_re_l
conv_6 (Conv2D) g2d_13[0][0]	(None, 24, 24, 256)	294912	max_poolin
norm_6 (BatchNormalization) 0]	(None, 24, 24, 256)	1024	conv_6[0][
leaky_re_lu_49 (LeakyReLU) 0]	(None, 24, 24, 256)	0	norm_6[0][
conv_7 (Conv2D) u_49[0][0]	(None, 24, 24, 128)	32768	leaky_re_l
norm_7 (BatchNormalization) 0]	(None, 24, 24, 128)	512	conv_7[0][
leaky_re_lu_50 (LeakyReLU) 0]	(None, 24, 24, 128)	0	norm_7[0][
conv_8 (Conv2D) u_50[0][0]	(None, 24, 24, 256)	294912	leaky_re_l
norm_8 (BatchNormalization) 0]	(None, 24, 24, 256)	1024	conv_8[0][
leaky_re_lu_51 (LeakyReLU) 0]	(None, 24, 24, 256)	0	norm_8[0][
max_pooling2d_14 (MaxPooling2D) u_51[0][0]	(None, 12, 12, 256)	0	leaky_re_l

conv_9 (Conv2D) g2d_14[0][0]	(None, 12, 12, 512)	1179648	max_poolin
norm_9 (BatchNormalization) 0]	(None, 12, 12, 512)	2048	conv_9[0][
leaky_re_lu_52 (LeakyReLU) 0]	(None, 12, 12, 512)	0	norm_9[0][
conv_10 (Conv2D) u_52[0][0]	(None, 12, 12, 256)	131072	leaky_re_l
norm_10 (BatchNormalization) [0]	(None, 12, 12, 256)	1024	conv_10[0]
leaky_re_lu_53 (LeakyReLU) [0]	(None, 12, 12, 256)	0	norm_10[0]
conv_11 (Conv2D) u_53[0][0]	(None, 12, 12, 512)	1179648	leaky_re_l
norm_11 (BatchNormalization) [0]	(None, 12, 12, 512)	2048	conv_11[0]
leaky_re_lu_54 (LeakyReLU) [0]	(None, 12, 12, 512)	0	norm_11[0]
conv_12 (Conv2D) u_54[0][0]	(None, 12, 12, 256)	131072	leaky_re_l
norm_12 (BatchNormalization) [0]	(None, 12, 12, 256)	1024	conv_12[0]
leaky_re_lu_55 (LeakyReLU) [0]	(None, 12, 12, 256)	0	norm_12[0]
conv_13 (Conv2D) u_55[0][0]	(None, 12, 12, 512)	1179648	leaky_re_l
norm_13 (BatchNormalization) [0]	(None, 12, 12, 512)	2048	conv_13[0]

leaky_re_lu_56 (LeakyReLU)	(None, 12, 12, 512)	0	norm_13[0]
max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 512)	0	leaky_re_lu_56[0][0]
conv_14 (Conv2D)	(None, 6, 6, 1024)	4718592	max_pooling2d_15[0][0]
norm_14 (BatchNormalization)	(None, 6, 6, 1024)	4096	conv_14[0]
leaky_re_lu_57 (LeakyReLU)	(None, 6, 6, 1024)	0	norm_14[0]
conv_15 (Conv2D)	(None, 6, 6, 512)	524288	leaky_re_lu_57[0][0]
norm_15 (BatchNormalization)	(None, 6, 6, 512)	2048	conv_15[0]
leaky_re_lu_58 (LeakyReLU)	(None, 6, 6, 512)	0	norm_15[0]
conv_16 (Conv2D)	(None, 6, 6, 1024)	4718592	leaky_re_lu_58[0][0]
norm_16 (BatchNormalization)	(None, 6, 6, 1024)	4096	conv_16[0]
leaky_re_lu_59 (LeakyReLU)	(None, 6, 6, 1024)	0	norm_16[0]
conv_17 (Conv2D)	(None, 6, 6, 512)	524288	leaky_re_lu_59[0][0]
norm_17 (BatchNormalization)	(None, 6, 6, 512)	2048	conv_17[0]
leaky_re_lu_60 (LeakyReLU)	(None, 6, 6, 512)	0	norm_17[0]

conv_18 (Conv2D) u_60[0][0]	(None, 6, 6, 1024)	4718592	leaky_re_l
norm_18 (BatchNormalization) [0]	(None, 6, 6, 1024)	4096	conv_18[0]
leaky_re_lu_61 (LeakyReLU) [0]	(None, 6, 6, 1024)	0	norm_18[0]
conv_19 (Conv2D) u_61[0][0]	(None, 6, 6, 1024)	9437184	leaky_re_l
norm_19 (BatchNormalization) [0]	(None, 6, 6, 1024)	4096	conv_19[0]
conv_21 (Conv2D) u_56[0][0]	(None, 12, 12, 64)	32768	leaky_re_l
leaky_re_lu_62 (LeakyReLU) [0]	(None, 6, 6, 1024)	0	norm_19[0]
norm_21 (BatchNormalization) [0]	(None, 12, 12, 64)	256	conv_21[0]
conv_20 (Conv2D) u_62[0][0]	(None, 6, 6, 1024)	9437184	leaky_re_l
leaky_re_lu_64 (LeakyReLU) [0]	(None, 12, 12, 64)	0	norm_21[0]
norm_20 (BatchNormalization) [0]	(None, 6, 6, 1024)	4096	conv_20[0]
lambda_2 (Lambda) u_64[0][0]	(None, 6, 6, 256)	0	leaky_re_l
leaky_re_lu_63 (LeakyReLU) [0]	(None, 6, 6, 1024)	0	norm_20[0]
concatenate_2 (Concatenate) [0]	(None, 6, 6, 1280)	0	lambda_2[0]

u_63[0][0]			leaky_re_l
conv_22 (Conv2D)	(None, 6, 6, 1024)	11796480	concatenate_2[0][0]
norm_22 (BatchNormalization)	(None, 6, 6, 1024)	4096	conv_22[0][0]
leaky_re_lu_65 (LeakyReLU)	(None, 6, 6, 1024)	0	norm_22[0][0]
flatten_2 (Flatten)	(None, 36864)	0	leaky_re_lu_65[0][0]
dropout_2 (Dropout)	(None, 36864)	0	flatten_2[0][0]
dense_2 (Dense)	(None, 4)	147460	dropout_2[0][0]
=====			
Total params: 50,695,396			
Trainable params: 50,674,724			
Non-trainable params: 20,672			